

# Data-Driven Models for Zebrafish Motion

## IDP final

---

Lukas Krenz

Advisers: Dr. Jacob Davidson (Constance), Nicola Rieke (CAMP)

Supervisor: Prof. Dr. Nassir Navab

July 1, 2018

TUM, Chair for Computer Aided Medical Procedures & Augmented Reality

Collaboration with Couzin Lab (Max Plank Institute for Ornithology/University of Constance)

# Introduction

Goal: Develop models for social behavior of juvenile zebrafish that extend to large groups.

Contributions:

- Comparison of models with increasing complexity for modeling interactions between two fish
- Data-driven spatial discretization
- Evaluation of importance of past trajectories
- Development and evaluation of a non-linear recurrent neural network that predicts parameters for mixture of Gaussians

Example **use case**: controlling a fish in a virtual reality environment  
“Pilot experiment” for neural network models for collective motion

## Zebrafish: Our Input Video

# Modeling Fish Motion

Data: Annotated (no tracking needed) videos from 10 experiments with 2 fish swimming, each for 1h. Annotations: positions of both fish and their orientation

Segmentation into c.a. 148000 kicks

Use wall model<sup>1</sup> to ignore areas with high wall influence (any wall closer than c.a. 4.8 cm).



Final data: c.a. 19400 kicks in training set (80% of all kicks)

We model kick trajectory (=kick direction and kick length).

---

<sup>1</sup>D. S. Calovi, A. Litchinko, V. Lecheval, U. Lopez, A. P. Escudero, H. Chaté, C. Sire and G. Theraulaz (2018). 'Disentangling and modeling interactions in fish with burst-and-coast swimming reveal distinct alignment and attraction behaviors'. In: *PLoS computational biology* 14.1, e1005933.

## Receptive Field<sup>2</sup>—Local coordinate system



Rotate coordinate system s.t. focal fish (red) has angle 0 (parallel to x-axis) and is at origin.

---

<sup>2</sup>R. Harpaz, G. Tkačik and E. Schneidman (2017). 'Discrete modes of social information processing predict individual behavior of fish in a group'. In: *Proceedings of the National Academy of Sciences*, p. 201703817.

## Receptive Field—Discretization

We discretize the area surrounding the fish in the local coordinate system.

We want:

- Symmetry around origin: to distinguish left/right, before/behind
- Bins should have equal number of fish.


Use  $8 \times 8$  bins.

Discretizing (with symmetry) axes independently is standard approach.

**But:** mean number of data points per bin of roughly  $2428 \pm 8483$  (mean  $\pm$  std.) for training and  $555 \pm 1926$  for testing.

12 bins **completely empty**, even for **training set**!


## Receptive Field—Data-driven discretization



Mean number of kicks per testing bin

Leads to  $2482 \pm 212$  and  $555 \pm 162$  data points per bin, instead of  $2428 \pm 8483$  and  $555 \pm 1926$ .

## Receptive Field—Data-driven discretization (zoomed in)



Mean number of kicks per testing bin




## Receptive Field—Extracted features

Bin number is one-hot encoded, alternative interpretation as 64 features, each is number of fish currently in bin.


Relative orientation of other fish encoded as unit vector with appropriate orientation, multiplied with bin-feature. Alternative interpretation as mean heading of fish in bin.

Extract this for timesteps 0 s, 0.05 s,  $\dots$ , 0.35 s before kick-off-time.



Overall  $64 \times (1 + 2) = 192$  features per timestep

# Social Models—Linear without memory



**No memory:** Only current position, etc.

$$\mathbf{y}^i = \mathbf{X}_0^i \beta_0^i + \text{bias} + \varepsilon$$

$$\varepsilon \sim \mathcal{N}(0, \sigma^2)$$


With:  $\mathcal{N}(\text{mean}, \text{variance})$  normal distribution

$\sigma$  standard deviation of residuals

$\mathbf{y}^i$  i-th component of output vector

$\mathbf{X}_0^i, \beta_0^i$  input matrix and weights for i-th component and timestep 0

# Social Models—Linear with memory



**Memory:** Current position and trace


**Concatenated:** Concatenate features for all timesteps then linear model

**Static:** Keep spatial weights  $\beta_t^i$  static for all  $t$

$$y^i = \sum_t c_t \mathbf{x}_t^i \beta^i + \varepsilon$$

with  $c_t$  mixing coefficient  
 $c_i \geq 0, \sum_i c_i = 1$ , normalized  
with *softmax*

# Social Models—Neural networks



General architecture. Symbols  $\approx$  and  $\mapsto$  correspond to encoding and decoding. Blue is input, gray hidden and black output.

Neural Networks can capture non-linear effects.


Encoder-decoder architecture with

**Encoder** transforms input features into hidden state

**Decoder** transforms hidden state into output

We discuss two encoders & two decoders.

# Neural Networks—Encoders: Multilayer-Perceptron




Architecture of MLP encoder. Symbol  $d(*)$  corresponds to linear layer, followed by tanh and dropout<sup>3</sup>. Blue is input, gray hidden state.

Multilayer perceptron encoder consists of two stacked layers (Linear  $\rightarrow$  Tanh  $\rightarrow$  Dropout) with 64 hidden neurons.

---

<sup>3</sup>N. Srivastava, G. Hinton, A. Krizhevsky, I. Sutskever and R. Salakhutdinov (2014). 'Dropout: A simple way to prevent neural networks from overfitting'. In: *The Journal of Machine Learning Research* 15.1, pp. 1929–1958.

# Neural Networks—Encoders: Recurrent neural network



Architecture of RNN encoder. Symbol  $*$  is linear layer,  $d()$  is recurrent dropout<sup>4</sup>,  $+$  is summation and  $f$  to  $\tanh$ . Blue is input, gray hidden state.

Input-to-hidden weights  $\mathbf{w}_{ih}$ , hidden-to-hidden weights  $\mathbf{w}_{hh}$ , biases  $\mathbf{b}$ , recurrent dropout  $d(\mathbf{x})$  (same mask for all timesteps), learned initial state of size 64 ( $h_0$ )

$$\mathbf{h}_t(\mathbf{X}) = \tanh(\mathbf{b} + \mathbf{w}_{ih}\mathbf{X}_t + \mathbf{w}_{hh}d(\mathbf{h}_{i-1})),$$

---

<sup>4</sup>S. Semeniuta, A. Severyn and E. Barth (2016). 'Recurrent dropout without memory loss'. In: *arXiv preprint arXiv:1603.05118*.

## Neural Networks—Decoder: Mixture density networks<sup>5</sup>

Predict mixture of Gaussians

$$p(\mathbf{y}|\mathbf{X}) = \sum_i^n \kappa_i \mathcal{N}(\boldsymbol{\mu}_i, \boldsymbol{\Sigma}_i),$$

with mixing coefficients  $\kappa_i$ , multivariate normal  $\mathcal{N}(\boldsymbol{\mu}_i, \boldsymbol{\Sigma}_i)$  with mean vector  $\boldsymbol{\mu}_i$  and covariance matrix  $\boldsymbol{\Sigma}_i$


Constraints:  $\kappa_i \geq 0$ ,  $\sum_i \kappa_i = 1$ , enforced by *softmax*

$\boldsymbol{\Sigma}_i$  valid (diagonal) covariance, diagonals (variances) must be positive, enforced by *exp*.

---

<sup>5</sup>C. M. Bishop (1994). *Mixture density networks*. Technical Report.

## Neural Networks—Decoder: Mixture density networks (cont.)<sup>6</sup>



Architecture of MDN decoder. Symbol  $*$  is linear layer. Gray is hidden state and black output.

---

<sup>6</sup>C. M. Bishop (1994). *Mixture density networks*. Technical Report.



## Results—Quantitative

Model	NLL-train	NLL-test	MSE-train	MSE-test
Baseline (Train-Mean)	2.25	2.01	0.558	0.423
MLP-MDN	1.63	<b>1.60</b>	0.448	0.373
RNN-MDN	<b>1.43</b>	1.69	0.436	<b>0.371</b>
MLP-MSE	2.04	1.87	0.452	0.375
RNN-MSE	2.00	1.86	<b>0.432</b>	0.377
Linear (w/o time)	2.04	1.87	0.451	0.376
Linear (time conc.)	2.00	1.86	0.432	0.373
Linear (static spatial)	2.04	1.86	0.451	0.374

Results for all models. Baseline: Always predict mean.

Note: Comparison of negative log likelihood NLL is valid, all models can be interpreted as mixture of Gaussians.

## Results—Example simulation with RNN-MDN

# Conclusion

- Data-driven discretization leads to more equal fish distribution than standard approach.
- Models extend trivially to larger fish groups.
- MDN-models allow sampling, multi-modal distributions and model uncertainty. This is **biologically more plausible** than our alternatives.
- Non-linear models work but do not show a significant improvement (but should work well for larger groups).